

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
# load dataset
df = pd.read_csv("creditcard.csv")
df.head(8)
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25	V26	V27	V28	Amount	Class
0	0.00	-1.36	-0.07	2.54	1.38	-0.34	0.46	0.24	0.10	0.36	...	-0.02	0.28	-0.11	0.07	0.13	-0.19	0.13	-0.02	149.62	0
1	0.00	1.19	0.27	0.17	0.45	0.06	-0.08	-0.08	0.09	-0.26	...	-0.23	-0.64	0.10	-0.34	0.17	0.13	-0.01	0.01	2.69	0
2	1.00	-1.36	-1.34	1.77	0.38	-0.50	1.80	0.79	0.25	-1.51	...	0.25	0.77	0.91	-0.69	-0.33	-0.14	-0.06	-0.06	378.66	0
3	1.00	-0.97	-0.19	1.79	-0.86	-0.01	1.25	0.24	0.38	-1.39	...	-0.11	0.01	-0.19	-1.18	0.65	-0.22	0.06	0.06	123.50	0
4	2.00	-1.16	0.88	1.55	0.40	-0.41	0.10	0.59	-0.27	0.82	...	-0.01	0.80	-0.14	0.14	-0.21	0.50	0.22	0.22	69.99	0
5	2.00	-0.43	0.96	1.14	-0.17	0.42	-0.03	0.48	0.26	-0.57	...	-0.21	-0.56	-0.03	-0.37	-0.23	0.11	0.25	0.08	3.67	0
6	4.00	1.23	0.14	0.05	1.20	0.19	0.27	-0.01	0.08	0.46	...	-0.17	-0.27	-0.15	-0.78	0.75	-0.26	0.03	0.01	4.99	0
7	7.00	-0.64	1.42	1.07	-0.49	0.95	0.43	1.12	-3.81	0.62	...	1.94	-1.02	0.06	-0.65	-0.42	-0.05	-1.21	-1.09	40.80	0

8 rows × 31 columns



```
# This shows the number of non fraudulent (0) and fraudulent (1) transactions
print(f"Unique values of target variable \n {df['Class'].unique()}")
print(f"Number of samples under each target value \n {df['Class'].value_counts()}")
```

```
Unique values of target variable
[0 1]
Number of samples under each target value
0    284315
1      492
Name: Class, dtype: int64
```

```
# time is irrelevant to classification so it can be removed
```

```
df = df.drop(['Time'], axis=1)
print(f"list of feature names after removing Time column :- \n{df.columns}")

list of feature names after removing Time column :-
Index(['V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10', 'V11',
      'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20', 'V21',
      'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount', 'Class'],
      dtype='object')

print(f"Dataset info : \n {df.info()}")
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 30 columns):
```

#	Column	Non-Null Count	Dtype
0	V1	284807 non-null	float64
1	V2	284807 non-null	float64
2	V3	284807 non-null	float64
3	V4	284807 non-null	float64
4	V5	284807 non-null	float64
5	V6	284807 non-null	float64
6	V7	284807 non-null	float64
7	V8	284807 non-null	float64
8	V9	284807 non-null	float64
9	V10	284807 non-null	float64
10	V11	284807 non-null	float64
11	V12	284807 non-null	float64
12	V13	284807 non-null	float64
13	V14	284807 non-null	float64
14	V15	284807 non-null	float64
15	V16	284807 non-null	float64
16	V17	284807 non-null	float64
17	V18	284807 non-null	float64
18	V19	284807 non-null	float64
19	V20	284807 non-null	float64
20	V21	284807 non-null	float64
21	V22	284807 non-null	float64
22	V23	284807 non-null	float64
23	V24	284807 non-null	float64
24	V25	284807 non-null	float64
25	V26	284807 non-null	float64
26	V27	284807 non-null	float64
27	V28	284807 non-null	float64
28	Amount	284807 non-null	float64
29	Class	284807 non-null	int64

```
dtypes: float64(29), int64(1)
memory usage: 65.2 MB
Dataset info :
None
```

```
from sklearn.preprocessing import StandardScaler
# amount needs a smaller range of numbers so it needs to be scaled.
df['norm_amount'] = StandardScaler().fit_transform(df['Amount'].values.reshape(-1,1))
fraud_df = df.drop(['Amount'], axis=1)
print(f"test values from Amount column after applying StandardScaler: \n {fraud_df['norm_amount'][0:4]}")
```

```
test values from Amount column after applying StandardScaler:
0    0.24
1   -0.34
2    1.16
3    0.14
Name: norm_amount, dtype: float64
```

```
# creating variables for train and test.
X = df.drop(['Class'], axis=1)
y = df[['Class']]
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(199364, 30)
(85443, 30)
(199364, 1)
(85443, 1)
```

The first implemented classifier is a neural network similar to assignment five

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```

y_train = np.asarray(y_train).astype('float32').reshape((-1,1))
y_test = np.asarray(y_test).astype('float32').reshape((-1,1))

from sklearn.model_selection import train_test_split
import tensorflow as tf
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Dropout
n_features = X_train.shape[1]

model = Sequential()
model.add(Dense(15,activation = 'relu',kernel_initializer = 'he_normal',input_shape = (n_features,)))
model.add(Dense(10,activation = 'relu',kernel_initializer = 'he_normal'))
model.add(Dropout(0.5))
model.add(Dense(1,activation='sigmoid'))

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy',tf.keras.metrics.AUC()])

model.fit(X_train, y_train,batch_size = 1000 ,epochs=150, verbose=2)

```

```

156/156 - 0s - loss: 0.0025 - accuracy: 0.9993 - auc: 0.9905 - 455ms/epoch - 3ms/step
Epoch 123/150
156/156 - 0s - loss: 0.0024 - accuracy: 0.9994 - auc: 0.9954 - 401ms/epoch - 3ms/step
Epoch 124/150
156/156 - 0s - loss: 0.0023 - accuracy: 0.9994 - auc: 0.9922 - 379ms/epoch - 2ms/step
Epoch 125/150
156/156 - 0s - loss: 0.0024 - accuracy: 0.9993 - auc: 0.9905 - 389ms/epoch - 2ms/step
Epoch 126/150
156/156 - 0s - loss: 0.0024 - accuracy: 0.9994 - auc: 0.9906 - 348ms/epoch - 2ms/step
Epoch 127/150
156/156 - 0s - loss: 0.0024 - accuracy: 0.9993 - auc: 0.9938 - 373ms/epoch - 2ms/step
Epoch 128/150
156/156 - 0s - loss: 0.0022 - accuracy: 0.9994 - auc: 0.9939 - 376ms/epoch - 2ms/step
Epoch 129/150
156/156 - 0s - loss: 0.0021 - accuracy: 0.9994 - auc: 0.9956 - 382ms/epoch - 2ms/step
Epoch 130/150
156/156 - 0s - loss: 0.0023 - accuracy: 0.9994 - auc: 0.9938 - 396ms/epoch - 3ms/step
Epoch 131/150
156/156 - 0s - loss: 0.0022 - accuracy: 0.9994 - auc: 0.9906 - 450ms/epoch - 3ms/step
Epoch 132/150
156/156 - 0s - loss: 0.0022 - accuracy: 0.9994 - auc: 0.9955 - 477ms/epoch - 3ms/step
Epoch 133/150
156/156 - 0s - loss: 0.0025 - accuracy: 0.9994 - auc: 0.9920 - 420ms/epoch - 3ms/step
Epoch 134/150

```

```

Epoch 134/150
156/156 - 0s - loss: 0.0024 - accuracy: 0.9993 - auc: 0.9905 - 443ms/epoch - 3ms/step
Epoch 135/150
156/156 - 0s - loss: 0.0024 - accuracy: 0.9993 - auc: 0.9921 - 411ms/epoch - 3ms/step
Epoch 136/150
156/156 - 0s - loss: 0.0022 - accuracy: 0.9994 - auc: 0.9922 - 416ms/epoch - 3ms/step
Epoch 137/150
156/156 - 0s - loss: 0.0026 - accuracy: 0.9993 - auc: 0.9919 - 437ms/epoch - 3ms/step
Epoch 138/150
156/156 - 0s - loss: 0.0021 - accuracy: 0.9995 - auc: 0.9956 - 405ms/epoch - 3ms/step
Epoch 139/150
156/156 - 0s - loss: 0.0023 - accuracy: 0.9994 - auc: 0.9905 - 404ms/epoch - 3ms/step
Epoch 140/150
156/156 - 0s - loss: 0.0022 - accuracy: 0.9994 - auc: 0.9906 - 356ms/epoch - 2ms/step
Epoch 141/150
156/156 - 0s - loss: 0.0024 - accuracy: 0.9993 - auc: 0.9904 - 364ms/epoch - 2ms/step
Epoch 142/150
156/156 - 0s - loss: 0.0025 - accuracy: 0.9994 - auc: 0.9970 - 421ms/epoch - 3ms/step
Epoch 143/150
156/156 - 0s - loss: 0.0023 - accuracy: 0.9994 - auc: 0.9939 - 370ms/epoch - 2ms/step
Epoch 144/150
156/156 - 0s - loss: 0.0022 - accuracy: 0.9994 - auc: 0.9956 - 390ms/epoch - 3ms/step
Epoch 145/150
156/156 - 0s - loss: 0.0022 - accuracy: 0.9994 - auc: 0.9939 - 402ms/epoch - 3ms/step
Epoch 146/150
156/156 - 0s - loss: 0.0025 - accuracy: 0.9993 - auc: 0.9938 - 397ms/epoch - 3ms/step
Epoch 147/150
156/156 - 0s - loss: 0.0022 - accuracy: 0.9994 - auc: 0.9957 - 404ms/epoch - 3ms/step
Epoch 148/150
156/156 - 0s - loss: 0.0023 - accuracy: 0.9994 - auc: 0.9956 - 387ms/epoch - 2ms/step
Epoch 149/150
156/156 - 0s - loss: 0.0022 - accuracy: 0.9994 - auc: 0.9956 - 403ms/epoch - 3ms/step
Epoch 150/150
156/156 - 0s - loss: 0.0022 - accuracy: 0.9994 - auc: 0.9955 - 402ms/epoch - 3ms/step
<keras.callbacks.History at 0x7fab1cb93d90>

```

The second classifier is the random forest

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report

def random_forest_classifier(X_train, y_train, X_test, y_test):
    # initialize object for DecisionTreeClassifier class

```

```

rf_classifier = RandomForestClassifier(n_estimators=50)
# train model by using fit method
print("Model training starts")
rf_classifier.fit(X_train, y_train.ravel())
acc_score = rf_classifier.score(X_test, y_test)
print(f'Accuracy of model on test dataset : {acc_score}')
# predict result using test dataset
y_pred = rf_classifier.predict(X_test)
# confusion matrix
print(f"Confusion Matrix : \n {confusion_matrix(y_test, y_pred)}")
# classification report for f1-score
print(f"Classification Report : \n {classification_report(y_test, y_pred)}")

```

```

# calling random_forest_classifier
random_forest_classifier(X_train, y_train, X_test, y_test)

```

```

Model training starts
Accuracy of model on test dataset : 0.9995903701883126
Confusion Matrix :
[[85298    9]
 [   26  110]]
Classification Report :

```

	precision	recall	f1-score	support
0.0	1.00	1.00	1.00	85307
1.0	0.92	0.81	0.86	136
accuracy			1.00	85443
macro avg	0.96	0.90	0.93	85443
weighted avg	1.00	1.00	1.00	85443

The third is the decision tree classifier

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report

def decision_tree_classification(X_train, y_train, X_test, y_test):
    # initialize object for DecisionTreeClassifier class
    dt_classifier = DecisionTreeClassifier()
    # train model by using fit method

```

```

print("Model training starts")
dt_classifier.fit(X_train, y_train.ravel())
print("Model training completed")
acc_score = dt_classifier.score(X_test, y_test)
print(f'Accuracy of model on test dataset : {acc_score}')
# predict result using test dataset
y_pred = dt_classifier.predict(X_test)
# confusion matrix
print(f"Confusion Matrix : \n {confusion_matrix(y_test, y_pred)}")
# classification report for f1-score
print(f"Classification Report : \n {classification_report(y_test, y_pred)}")

```

```

# calling decision_tree_classification method to train and evaluate model
decision_tree_classification(X_train, y_train, X_test, y_test)

```

```

Model training starts
Model training completed
Accuracy of model on test dataset : 0.999133925541004
Confusion Matrix :
[[85263   44]
 [   30  106]]
Classification Report :

```

	precision	recall	f1-score	support
0.0	1.00	1.00	1.00	85307
1.0	0.71	0.78	0.74	136
accuracy			1.00	85443
macro avg	0.85	0.89	0.87	85443
weighted avg	1.00	1.00	1.00	85443

The beginning of the Regression section of the project.

```

import pandas as pd
import pandas_profiling
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
sns.set_style('whitegrid')

```


```
df2 = pd.read_excel('ENB2012_data.xlsx')
```

```
df2.columns=["relative_compactness","surface_area","wall_area","roof_area","overall_height","orientation",  
            "glazing_area","glazing_area_dist","heating_load","cooling_load"]
```

df2











	relative_compactness	surface_area	wall_area	roof_area	overall_height	orientation	glazing_area	glazing_area_dist	hea
0	0.98	514.50	294.00	110.25	7.00	2	0.00	0	
1	0.98	514.50	294.00	110.25	7.00	3	0.00	0	
2	0.98	514.50	294.00	110.25	7.00	4	0.00	0	
3	0.98	514.50	294.00	110.25	7.00	5	0.00	0	
4	0.90	563.50	318.50	122.50	7.00	2	0.00	0	
...	...	...	...	...	...	...	...	...	...
763	0.64	784.00	343.00	220.50	3.50	5	0.40	5	
764	0.62	808.50	367.50	220.50	3.50	2	0.40	5	
765	0.62	808.50	367.50	220.50	3.50	3	0.40	5	
766	0.62	808.50	367.50	220.50	3.50	4	0.40	5	
767	0.62	808.50	367.50	220.50	3.50	5	0.40	5	

768 rows × 10 columns



```
df2.hist(figsize=(15,15))  
plt.show()
```






```

fig, axs = plt.subplots(1, 8, sharey=True)
df2.plot(kind='scatter', x='relative_compactness', y='heating_load', ax=axs[0], figsize=(14, 6))
df2.plot(kind='scatter', x='surface_area', y='heating_load', ax=axs[1])
df2.plot(kind='scatter', x='wall_area', y='heating_load', ax=axs[2])
df2.plot(kind='scatter', x='roof_area', y='heating_load', ax=axs[3])
df2.plot(kind='scatter', x='overall_height', y='heating_load', ax=axs[4])
df2.plot(kind='scatter', x='orientation', y='heating_load', ax=axs[5])
df2.plot(kind='scatter', x='glazing_area', y='heating_load', ax=axs[6])
df2.plot(kind='scatter', x='glazing_area_dist', y='heating_load', ax=axs[7])

```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fab9f82a4d0>



```
df2.isnull().sum()
```

```

relative_compactness    0
surface_area            0
wall_area               0
roof_area               0
overall_height          0
orientation              0
glazing_area            0
glazing_area_dist       0
heating_load            0

```

```
cooling_load      0
dtype: int64
```

df2.info

```
<bound method DataFrame.info of
0      0.98      514.5      294.0      110.25      7.0
1      0.98      514.5      294.0      110.25      7.0
2      0.98      514.5      294.0      110.25      7.0
3      0.98      514.5      294.0      110.25      7.0
4      0.90      563.5      318.5      122.50      7.0
..      ...      ...      ...      ...      ...
763     0.64      784.0      343.0      220.50      3.5
764     0.62      808.5      367.5      220.50      3.5
765     0.62      808.5      367.5      220.50      3.5
766     0.62      808.5      367.5      220.50      3.5
767     0.62      808.5      367.5      220.50      3.5

orientation  glazing_area  glazing_area_dist  heating_load  cooling_load
0             2           0.0                0           15.55        21.33
1             3           0.0                0           15.55        21.33
2             4           0.0                0           15.55        21.33
3             5           0.0                0           15.55        21.33
4             2           0.0                0           20.84        28.28
..      ...      ...      ...      ...      ...
763           5           0.4                5           17.88        21.40
764           2           0.4                5           16.54        16.88
765           3           0.4                5           16.44        17.11
766           4           0.4                5           16.48        16.61
767           5           0.4                5           16.64        16.03
```

```
[768 rows x 10 columns]>
```

df2.describe

```
<bound method NDFrame.describe of
0      0.98      514.5      294.0      110.25      7.0
1      0.98      514.5      294.0      110.25      7.0
2      0.98      514.5      294.0      110.25      7.0
3      0.98      514.5      294.0      110.25      7.0
4      0.90      563.5      318.5      122.50      7.0
..      ...      ...      ...      ...      ...
763     0.64      784.0      343.0      220.50      3.5
764     0.62      808.5      367.5      220.50      3.5
```


765	0.62	808.5	367.5	220.50	3.5
766	0.62	808.5	367.5	220.50	3.5
767	0.62	808.5	367.5	220.50	3.5

	orientation	glazing_area	glazing_area_dist	heating_load	cooling_load
0	2	0.0	0	15.55	21.33
1	3	0.0	0	15.55	21.33
2	4	0.0	0	15.55	21.33
3	5	0.0	0	15.55	21.33
4	2	0.0	0	20.84	28.28
..	...	...	...	...	...
763	5	0.4	5	17.88	21.40
764	2	0.4	5	16.54	16.88
765	3	0.4	5	16.44	17.11
766	4	0.4	5	16.48	16.61
767	5	0.4	5	16.64	16.03

[768 rows x 10 columns]>

```
plt.figure(figsize=(12,12))
sns.heatmap(df2.corr(),annot=True)
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7faba034d610>



```
pd.set_option('display.float_format',lambda x: '{:,.2f}'.format(x) if abs(x) < 10000 else '{:,.0f}'.format(x))
df2.corr()
```

	relative_compactness	surface_area	wall_area	roof_area	overall_height	orientation	glazing_area	glazin
<b>relative_compactness</b>	1.00	-0.99	-0.20	-0.87	0.83	0.00	-0.00	
<b>surface area</b>	-0.99	1.00	0.20	0.88	-0.86	-0.00	0.00	


```
#Normalize inputs, set output
from sklearn.preprocessing import Normalizer
nr = Normalizer(copy=False)

X = df2.drop(['heating_load','cooling_load'], axis=1)
X = nr.fit_transform(X)
y = df2[['heating_load','cooling_load']]
```

.....

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 123)
```

```
X_new = pd.DataFrame({'relative_compactness': [df2.relative_compactness.min(), df2.relative_compactness.max()]})
X_new.head()
```

	relative_compactness	
0	0.62	
1	0.98	

```
import statsmodels.formula.api as smf

# Create fitted model.
lm = smf.ols(formula='heating_load ~ relative_compactness', data=df2).fit()


# Get coefficients.
lm.params

pred = lm.predict(X_new)
pred

0    13.75
1    35.12
dtype: float64
```

```
df2.plot(kind='scatter', x='relative_compactness', y='heating_load')
plt.plot(X_new, pred, c='red', linewidth=2)
```

[<matplotlib.lines.Line2D at 0x7faba03bd590>]



## Implementation of Lasso Regression

```
from sklearn.linear_model import Lasso
from sklearn.linear_model import LassoCV
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
```

```
lasso_model = Lasso().fit(X_train, y_train)
lasso_model
```

```
Lasso()
```

```
lasso_model.intercept_
```

```
array([22.101875 , 24.35225694])
```

```
lasso_model.coef_
```

```
array([[ 0., -0.,  0., -0.,  0.,  0.,  0.,  0.]
```

```
[ 0., -0., 0., -0., 0., 0., 0., 0.]])
```

```
from sklearn.metrics import mean_squared_error
```

```
y_pred = lasso_model.predict(X_test)
np.sqrt(mean_squared_error(y_test, y_pred))
```

```
10.000941383873894
```

```
lasso_model.score(X,y)
```

```
-0.0005141088871829513
```

## Implementation of Ridge Regression

```
from sklearn.linear_model import Ridge
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = 42) # Train test split
ridge_model = Ridge(alpha = 5).fit(X_train, y_train)
ridge_model
```

```
Ridge(alpha=5)
```

```
ridge_model.intercept_
```

```
array([22.32902789, 24.52907455])
```

```
ridge_model.coef_
```

```
array([[ 0.11629168, -11.54544662, 36.83527657, -24.19036159,
         1.88144148,  0.24205811,  0.07361704,  0.41832675],
       [ 0.11410307, -10.71744662, 34.30447853, -22.51096257,
         1.81488642,  0.2554495 ,  0.05716537,  0.31520485]])
```

```
ridge_model = Ridge().fit(X_train, y_train)
y_pred = ridge_model.predict(X_train)
RMSE = np.sqrt(mean_squared_error(y_train, y_pred))
RMSE
```

```
5.314821337489512
```



```
y_pred = ridge_model.predict(X_test)
RMSE = np.sqrt(mean_squared_error(y_test, y_pred))
RMSE
```

```
5.328363319249427
```

```
ridge_model.score(X,y)
```

```
0.7049581597106621
```

✓ 21s completed at 6:21 PM

